

```

import os
import datetime
from datetime import timedelta
import urllib.request
from matplotlib.path import Path
from matplotlib.pyplot import fill_between
from pyparsing import withClass
from pyspark import SparkConf
from pyspark import SparkContext
from pyspark.sql import SparkSession, Window
import sys
from collections import namedtuple
from pyspark.sql.functions import *
from pyspark.sql.types import *
from pyspark.sql.functions import *
import builtins
import random
import urllib.request

python_path = sys.executable
os.environ['PYSPARK_PYTHON'] = python_path
os.environ['JAVA_HOME'] = r'C:\Users\LENOVO\.jdk\corretto-1.8.0_422'
os.environ['HADOOP_HOME'] = "C:\Program Files\Hadoop"

conf = (SparkConf()
        .setAppName("pyspark")
        .setMaster("local[*]")
        .set("spark.driver.host", "localhost")
        .set("spark.default.parallelism", "1")
        )
sc = SparkContext(conf=conf)
spark = SparkSession.builder.getOrCreate()

# SPARK BOILERPLATE
# """
#
#
# """
# # input_df1
# +-----+-----+-----+-----+-----+-----+
# |emp_id|  name|department_id| hire_date|salary|age|
# +-----+-----+-----+-----+-----+-----+
# |    1| Alice|          101|2020-01-15| 70000| 29|
# |    2|   Bob|          102|2018-06-23| 60000| 35|
# |    3|Charlie|          101|2019-07-01| 90000| 42|
# |    4| David|          103|2021-03-12| 50000| 24|
# |    5|   Eve|          102|2020-09-10| 75000| 30|
# +-----+-----+-----+-----+-----+-----+
# # input_df2
# +-----+-----+-----+
# |department_id|department_name|  location|
# +-----+-----+-----+
# |          101|   Engineering|   New York|
# |          102|   Marketing|San Francisco|
# |          103|     Sales|    Chicago|
# +-----+-----+-----+
# # result_df
# #note: find the average salary and the max_age by department
# +-----+-----+-----+
# |department_name|avg_salary|max_age|
# +-----+-----+-----+
# |   Engineering| 80000.0|   42|
# |   Marketing| 67500.0|   35|
# |     Sales| 50000.0|   24|
# +-----+-----+-----+
#
#
# """

```

```

# ## Sample DataFrames
employees = [(1, 'Alice', 101, '2020-01-15', 70000, 29),
             (2, 'Bob', 102, '2018-06-23', 60000, 35),
             (3, 'Charlie', 101, '2019-07-01', 90000, 42),
             (4, 'David', 103, '2021-03-12', 50000, 24),
             (5, 'Eve', 102, '2020-09-10', 75000, 30)]

departments = [(101, 'Engineering', 'New York'),
               (102, 'Marketing', 'San Francisco'),
               (103, 'Sales', 'Chicago')]

# Creating DataFrames
df_employees = spark.createDataFrame(
    employees,
    ['emp_id', 'name', 'department_id', 'hire_date', 'salary', 'age']
)

df_departments = spark.createDataFrame(departments, ['department_id', 'department_name', 'location'])

joined_df = (df_departments
             .join(df_employees, "department_id", "left")
             .groupBy("department_name")
             .agg(
                 avg("salary").alias("avg_salary"),
                 max("age").alias("max_age")
             ).orderBy("department_name")) #.select("department_name", "avg_salary", "max_age")

joined_df.show()
#note: with spark.sql
df_employees.createOrReplaceTempView("employees")
df_departments.createOrReplaceTempView("departments")
spark.sql("""
    select
        d.department_name, avg(e.salary), max(e.age)
    from departments d
        left join employees e
            on d.department_id = e.department_id
    group by department_name
    order by department_name
""").show()

```

```

#
# # info: scenario 2
# """
# # Convert price from string to double.
# # Group by product_name and aggregate the total sales (quantity * price).
# # Separate aggregation based on customer_type: calculate total sales for
# # Regular and Premium customers.
# # product_name    total_sales_regular    total_sales_premium
# +-----+-----+-----+-----+-----+-----+
# |transaction_id|product_name|quantity| price| sale_date|customer_type|
# +-----+-----+-----+-----+-----+-----+
# |          1|      Laptop|         1|1200.5|2023-01-10|      Regular|
# |          2|       Mouse|         3|  25.0|2023-01-11|      Regular|
# |          3|      Laptop|         2|1200.5|2023-01-12|     Premium|
# |          4|   Keyboard|         1|  75.0|2023-01-13|      Regular|
# |          5|       Mouse|         5|  25.0|2023-01-14|     Premium|
# +-----+-----+-----+-----+-----+-----+
#
# +-----+-----+-----+
# |product_name|total_sales_regular|total_sales_Premium|
# +-----+-----+-----+
# |      Laptop|          1200.5|          2401.0|
# |       Mouse|           75.0|           125.0|
# |   Keyboard|           75.0|              0.0|
# +-----+-----+-----+

```

```

transactions = [(1, 'Laptop', 1, '1200.50', '2023-01-10', 'Regular'),
                (2, 'Mouse', 3, '25.00', '2023-01-11', 'Regular'),
                (3, 'Laptop', 2, '1200.50', '2023-01-12', 'Premium'),
                (4, 'Keyboard', 1, '75.00', '2023-01-13', 'Regular'),
                (5, 'Mouse', 5, '25.00', '2023-01-14', 'Premium')]

# Create DataFrame
df_transactions = (spark
                  .createDataFrame(
                      transactions,
                      ['transaction_id', 'product_name', 'quantity', 'price', 'sale_date',
                      'customer_type']))
df_transactions = df_transactions.withColumn('price', col('price').cast('double'))
df_transactions.show()

```

```

sales_groupby_df = (df_transactions.groupby("product_name").agg(
    sum(
        when(col("customer_type") == 'Regular', col("quantity") * col("price"))
        .otherwise(0))
    .alias("total_sales_regular"),
    sum(
        when(col("customer_type") == 'Premium', col("quantity") * col("price"))
        .otherwise(0)).alias("total_sales_Premium")
))
sales_groupby_df.show()

```

```
print(df_transactions.dtypes)
```

```
df_transactions.createOrReplaceTempView("transactions")
```

```

df = spark.sql("""
select
    product_name,
    sum(case
        when customer_type='Regular'
        then quantity * cast(price as double) else 0 end) as Regular_sales,
    sum(case
        when customer_type='Premium'
        then quantity * cast(price as double) else 0 end) as Premium_sales
from
    transactions
group by
    product_name
""")
df.show()
print(df.dtypes)

```

```

#
# #note: scenario-3
# # Join sales with products on product_id.
# # Handle NULL values by replacing NULL in quantity with 1 and NULL in price with 100.
# # Group by category and aggregate the total sales and count of distinct stores.
# # category,    total_sales,    distinct_stores
# """

```

```
# Two input DF's
```

```

# +-----+-----+-----+-----+
# |product_id|product_name|    category|manufacturer|
# +-----+-----+-----+-----+
# |      1001|      Phone|Electronics|    XYZ Corp|
# |      1002|    Charger|Electronics|    ABC Ltd|
# |      1003|Headphones|Electronics|    XYZ Corp|
# |      1004|Phone Cover|Accessories|    PQR Inc|
# |      1005|    Tablet|Electronics|    XYZ Corp|
# +-----+-----+-----+-----+
#

```

```

# +-----+-----+-----+-----+-----+
# |sale_id|product_id|quantity|price| sale_date|store_id|
# +-----+-----+-----+-----+-----+
# |      1|      1001|      2|500.0|2023-08-01|      201|
# |      2|      1002|     NULL| 20.0|2023-08-02|     NULL|
# |      3|      1003|      4| NULL|2023-08-03|      202|
# |      4|      1004|      1|300.0|2023-08-04|      201|
# |      5|      1005|     NULL| NULL|2023-08-05|      203|
# +-----+-----+-----+-----+-----+
#
# the Resultant df
# +-----+-----+-----+
# | category|total_sales|distinct_stores|
# +-----+-----+-----+
# |Electronics|      1520.0|              3|
# |Accessories|      300.0|              1|
# +-----+-----+-----+
# """
sales = [(1, 1001, 2, 500.00, '2023-08-01', 201),
         (2, 1002, None, 20.00, '2023-08-02', None),
         (3, 1003, 4, None, '2023-08-03', 202),
         (4, 1004, 1, 300.00, '2023-08-04', 201),
         (5, 1005, None, None, '2023-08-05', 203)]

products = [(1001, 'Phone', 'Electronics', 'XYZ Corp'),
            (1002, 'Charger', 'Electronics', 'ABC Ltd'),
            (1003, 'Headphones', 'Electronics', 'XYZ Corp'),
            (1004, 'Phone Cover', 'Accessories', 'PQR Inc'),
            (1005, 'Tablet', 'Electronics', 'XYZ Corp')]

# Create DataFrames
df_sales = spark.createDataFrame(
    sales,
    ['sale_id', 'product_id', 'quantity', 'price', 'sale_date', 'store_id']
)
df_products = spark.createDataFrame(
    products,
    ['product_id', 'product_name', 'category', 'manufacturer']
)

df_products.show()
df_sales.show()

joined_df = df_sales.join(df_products, "product_id", "left")
null_df = joined_df.fillna({"quantity": 1, "price": 100})
null_df.show()
res1 = (null_df
        .groupby("category")
        .agg(
            sum(col("price") * col("quantity")).alias("total_sales"),
            countDistinct("store_id").alias("distinct_stores")
        )
    )
res1.show()

df_products.createOrReplaceTempView("products")
df_sales.createOrReplaceTempView("sales")

```

```
spark.sql("""
select
    p.category,
    sum(coalesce(s.quantity, 1)* coalesce(s.price,100)) as total_sales,
    count(distinct store_id) as distinct_stores
from sales s
    left join products p
on s.product_id = p.product_id
group by p.category
order by distinct_stores desc
""").show()

# # note: scenario 4
# # Apply window functions to calculate the cumulative sum of page_views for each user_id,
# # ordered by visit_time.
#
# # Calculate the difference between the current page_views and the previous value.
# """
# i/p
# +-----+-----+-----+-----+
# |visit_id|user_id|page_views|          visit_time|
# +-----+-----+-----+-----+
# |      1|    101|         5|2024-01-10 09:00:00|
# |      2|    101|         7|2024-01-10 10:00:00|
# |      3|    102|         3|2024-01-10 09:30:00|
# |      4|    101|         4|2024-01-10 11:00:00|
# |      5|    103|        10|2024-01-10 10:30:00|
# +-----+-----+-----+-----+
# o/p
# +-----+-----+-----+-----+-----+-----+
# |visit_id|user_id|page_views|          visit_time|cum_sum_pages|diff_from_last_view|
# +-----+-----+-----+-----+-----+-----+
# |      1|    101|         5|2024-01-10 09:00:00|          5|          NULL|
# |      2|    101|         7|2024-01-10 10:00:00|         12|             2|
# |      4|    101|         4|2024-01-10 11:00:00|         16|            -3|
# |      3|    102|         3|2024-01-10 09:30:00|          3|          NULL|
# |      5|    103|        10|2024-01-10 10:30:00|         10|          NULL|
# +-----+-----+-----+-----+-----+-----+
# """
#
# # # Sample DataFrame
web_traffic = [(1, 101, 5, '2024-01-10 09:00:00'),
               (2, 101, 7, '2024-01-10 10:00:00'),
               (3, 102, 3, '2024-01-10 09:30:00'),
               (4, 101, 4, '2024-01-10 11:00:00'),
               (5, 103, 10, '2024-01-10 10:30:00')]

# Create DataFrame
df_web_traffic = spark.createDataFrame(
    web_traffic,
    ['visit_id', 'user_id', 'page_views', 'visit_time']
)
df_web_traffic.show()

window_spec = Window.partitionBy("user_id").orderBy("visit_time")
res = (df_web_traffic
      .withColumn("cum_sum_pages",sum("page_views").over(window_spec))
      .withColumn(
          "diff_from_last_view",
          col("page_views") - lag("page_views").over(window_spec)
      )
    )
```

```
res.show()

df_web_traffic.createOrReplaceTempView("traffic")
spark.sql("""
select
    visit_id, user_id, visit_time, page_views,
    sum(page_views) over(partition by user_id order by visit_time) as cum_sum_pages,
    page_views-lag(page_views) over( partition by user_id order by visit_time) as diff_time
from traffic
""").show()
```

```
# # note: scenario 5
# # You want to assign ranks to sales representatives based on their total sales,
# # partitioned by department.
```

```
# """
```

```
# i/p
```

```
# +-----+-----+-----+-----+
# |sale_id|department|rep_name|sales_amount|
# +-----+-----+-----+-----+
# |      1|      Sales|    Alice|         200|
# |      2|      Sales|     Bob|         150|
# |      3| Marketing|  Charlie|         300|
# |      4| Marketing|    David|         100|
# |      5|      Sales|     Eve|         250|
# +-----+-----+-----+-----+
```

```
#
```

```
# o/p
```

```
# +-----+-----+-----+-----+-----+-----+
# |sale_id|department|rep_name|sales_amount|rank|row_number|
# +-----+-----+-----+-----+-----+-----+
# |      3| Marketing|  Charlie|         300|    1|         1|
# |      4| Marketing|    David|         100|    2|         2|
# |      5|      Sales|     Eve|         250|    1|         1|
# |      1|      Sales|    Alice|         200|    2|         2|
# |      2|      Sales|     Bob|         150|    3|         3|
# +-----+-----+-----+-----+-----+-----+
```

```
# """
```

```
# # Sample DataFrame
```

```
sales_data = [(1, 'Sales', 'Alice', 200),
              (2, 'Sales', 'Bob', 150),
              (3, 'Marketing', 'Charlie', 300),
              (4, 'Marketing', 'David', 100),
              (5, 'Sales', 'Eve', 250)]
```

```
df_sales = spark.createDataFrame(sales_data, ['sale_id', 'department', 'rep_name', 'sales_amount'])
```

```
df_sales.show()
```

```
window_spec = Window.partitionBy("department").orderBy(desc("sales_amount"))
```

```
res = (df_sales
      .withColumn("rank", rank().over(window_spec))
      .withColumn("row_number", row_number().over(window_spec))
      )
res.show()
```

```
df_sales.createOrReplaceTempView("sales")
```

```
spark.sql("""
select
    sale_id, department, rep_name, sales_amount,
    rank()over(partition by department order by sales_amount desc) as rank,
    dense_rank()over(partition by department order by sales_amount desc) as dense_rank,
    row_number()over(partition by department order by sales_amount desc) as row_num
from sales
""").show()
```

```

# #note: scenario 6
# # Calculate a moving average of daily sales over a 3-day window.
# """
# i/p
# +-----+-----+-----+
# |sale_id|      date|sales_amount|
# +-----+-----+-----+
# |      1|2024-01-01|          100|
# |      2|2024-01-02|          150|
# |      3|2024-01-03|          200|
# |      4|2024-01-04|          250|
# |      5|2024-01-05|          300|
# +-----+-----+-----+
#
# o/p
# +-----+-----+-----+-----+
# |sale_id|      date|sales_amount|moving_avg|
# +-----+-----+-----+-----+
# |      1|2024-01-01|          100|      100.0|
# |      2|2024-01-02|          150|      150.0|
# |      3|2024-01-03|          200|      200.0|
# |      4|2024-01-04|          250|      250.0|
# |      5|2024-01-05|          300|      300.0|
# +-----+-----+-----+-----+
# """
daily_sales = [(1, '2024-01-01', 100),
               (2, '2024-01-02', 150),
               (3, '2024-01-03', 200),
               (4, '2024-01-04', 250),
               (5, '2024-01-05', 300)]

# # Calculate a moving average of daily sales over a 3-day window.
df_daily_sales = spark.createDataFrame(daily_sales, ['sale_id', 'date', 'sales_amount'])

df_daily_sales.show()
window_spec = Window.orderBy("date").rowsBetween(-2,0)
res = df_daily_sales.withColumn("moving_avg", avg('sales_amount').over(window_spec))
res.show()

df_daily_sales.createOrReplaceTempView("sales")
spark.sql("""
    select sale_id, date, sales_amount,
    avg(sales_amount) over(order by date rows between 2 preceding and current row) as moving_avg
    from sales
""").show()

# #note: scenario 7
# # Calculate the percentile rank of each employee's salary within their department.
# """
# i/p
# +-----+-----+-----+-----+
# |emp_id|department|emp_name|salary|
# +-----+-----+-----+-----+
# |      1|      IT|   Alice| 90000|
# |      2|      IT|    Bob| 75000|
# |      3|     HR| Charlie| 70000|
# |      4|     HR|   David| 80000|
# |      5|      IT|    Eve| 95000|
# |      6|      IT|    Eve|  5000|
# +-----+-----+-----+-----+
#

```

```
# o/p
# +-----+-----+-----+-----+
# |emp_id|department|emp_name|salary|rank|
# +-----+-----+-----+-----+
# |      4|      HR|   David| 80000|  0.0|
# |      3|      HR| Charlie| 70000|100.0|
# |      5|      IT|    Eve| 95000|  0.0|
# |      1|      IT|   Alice| 90000|33.33333333333333|
# |      2|      IT|    Bob| 75000|66.66666666666666|
# |      6|      IT|    Eve| 5000|100.0|
# +-----+-----+-----+-----+
# """
employee_salaries = [(1, 'IT', 'Alice', 90000),
                     (2, 'IT', 'Bob', 75000),
                     (3, 'HR', 'Charlie', 70000),
                     (4, 'HR', 'David', 80000),
                     (5, 'IT', 'Eve', 95000),
                     (6, 'IT', 'Eve', 5000)]

df_salaries = spark.createDataFrame(
    employee_salaries,
    ['emp_id', 'department', 'emp_name', 'salary']
)
df_salaries.show()

window_spec = Window.partitionBy("department").orderBy(desc("salary"))
res = df_salaries.withColumn("rank", percent_rank().over(window_spec)*100)
res.show()

df_with_percentile_rank = df_salaries.withColumn(
    'percentile_rank',
    ((rank().over(window_spec) - 1) / (count('emp_id').over(Window.partitionBy('department')) - 1)) * 100
)
df_with_percentile_rank.show(truncate=8)

df_salaries.createOrReplaceTempView("salaries")
spark.sql("""
    select
        emp_id, department, emp_name, salary,
        percent_rank()over(partition by department order by salary desc) as percentile_rank,

        ((RANK() OVER (PARTITION BY department ORDER BY salary) - 1) * 100.0 /
        (COUNT(emp_id) OVER (PARTITION BY department) - 1)) AS math_percentile_rank
    from salaries
""").show()

# # note: scenario 8
# # Compute the difference in sales between the current and previous day.
# """
# i/p
# +-----+-----+-----+
# |sale_id|      date|sales_amount|
# +-----+-----+-----+
# |      1|2024-01-01|          100|
# |      2|2024-01-02|          150|
# |      3|2024-01-03|          200|
# |      4|2024-01-04|          250|
# |      5|2024-01-05|          300|
# +-----+-----+-----+
#
```



```

# o/p
# +-----+-----+-----+-----+-----+
# |sale_id|      date|sales_amount|prev_day_sales|next_day_sales|sales_diff|
# +-----+-----+-----+-----+-----+
# |      1|2024-01-01|          100|          NULL|          150|          NULL|
# |      2|2024-01-02|          150|          100|          200|           50|
# |      3|2024-01-03|          200|          150|          250|           50|
# |      4|2024-01-04|          250|          200|          300|           50|
# |      5|2024-01-05|          300|          250|          NULL|           50|
# +-----+-----+-----+-----+-----+
# """
daily_sales = [(1, '2024-01-01', 100),
               (2, '2024-01-02', 150),
               (3, '2024-01-03', 200),
               (4, '2024-01-04', 350),
               (5, '2024-01-05', 300)]

df_daily_sales = spark.createDataFrame(daily_sales, ['sale_id', 'date', 'sales_amount'])
df_daily_sales.show()
window_spec = Window.orderBy("date")

res = df_daily_sales.withColumn(
    "prev_day_sales", lag("sales_amount").over(window_spec)
).withColumn(
    "next_day_sales", lead("sales_amount").over(window_spec)
).withColumn("sales_diff", col("sales_amount") - col("prev_day_sales"))
res.show()

df_daily_sales.createOrReplaceTempView("sales")
spark.sql("""
    select
        sale_id, date, sales_amount,
        lag(sales_amount) over(order by date) as prev_day_sales,
        lead(sales_amount) over(order by date) as next_day_sales,
        (sales_amount - prev_day_sales) as sales_diff
    from sales
""").show()

```

```

# #note: scenario 9
# """
# i/p
# +-----+-----+-----+-----+
# |sale_id|rep_name|      date|sales_amount|
# +-----+-----+-----+-----+
# |      1|  Alice|2024-01-01|          100|
# |      2|  Alice|2024-01-02|          150|
# |      3|   Bob|2024-01-01|          200|
# |      4|   Bob|2024-01-02|          250|
# +-----+-----+-----+-----+
# o/p
# +-----+-----+-----+-----+
# |sale_id|rep_name|      date|sales_amount|running_total|
# +-----+-----+-----+-----+
# |      1|  Alice|2024-01-01|          100|          100|
# |      2|  Alice|2024-01-02|          150|          250|
# |      3|   Bob|2024-01-01|          200|          200|
# |      4|   Bob|2024-01-02|          250|          450|
# +-----+-----+-----+-----+
# """

```

```
sales_data = [(1, 'Alice', '2024-01-01', 100),
              (2, 'Alice', '2024-01-02', 150),
              (3, 'Bob', '2024-01-01', 200),
              (4, 'Bob', '2024-01-02', 250)]

df_sales = spark.createDataFrame(sales_data, ['sale_id', 'rep_name', 'date', 'sales_amount'])
df_sales.show()

# Define window specification
window_spec = (Window
              .partitionBy('rep_name')
              .orderBy('date')
              .rowsBetween(Window.unboundedPreceding, Window.currentRow)
              )

# Calculate running total
df_with_running_total = (
    df_sales.withColumn('running_total', sum('sales_amount').over(window_spec))
)

df_with_running_total.show()

df_sales.createOrReplaceTempView("sales")
spark.sql("""
    select
        sale_id, rep_name, date, sales_amount,
        sum(sales_amount) over(partition by rep_name order by date) as running_total
    from sales
""")
```

*# #note: Scenario 10: Calculate the difference in sales between the current day and the next day for each representative.*

```
# """
# i/p
# +-----+-----+-----+-----+
# |sale_id|rep_name|    date|sales_amount|
# +-----+-----+-----+-----+
# |      1|  Alice|2024-01-01|      100|
# |      2|  Alice|2024-01-02|      150|
# |      3|   Bob|2024-01-01|      200|
# |      4|   Bob|2024-01-02|      250|
# +-----+-----+-----+-----+
# """
```

*# o/p*

sale_id	rep_name	date	sales_amount	next_day_sales	prev_day_sales	sales_diff_next	sales_diff_prev
1	Alice	2024-01-01	100	150	NULL	50	NULL
2	Alice	2024-01-02	150	NULL	100	NULL	50
3	Bob	2024-01-01	200	250	NULL	50	NULL
4	Bob	2024-01-02	250	NULL	200	NULL	50

```
"""
# """
#
# # note:
```

```
sales_data = [(1, 'Alice', '2024-01-01', 100),
              (2, 'Alice', '2024-01-02', 150),
              (3, 'Bob', '2024-01-01', 200),
              (4, 'Bob', '2024-01-02', 250)]

df_sales = spark.createDataFrame(sales_data, ['sale_id', 'rep_name', 'date', 'sales_amount'])
df_sales.show()

# Define window specification
window_spec = Window.partitionBy('rep_name').orderBy('date')
res = ((df_sales
        .withColumn("next_day_sales", lead("sales_amount").over(window_spec)))
        .withColumn("prev_day_sales", lag("sales_amount").over(window_spec)))
        .withColumn("sales_diff_next", col("next_day_sales")-col("sales_amount"))
        .withColumn("sales_diff_prev", col("sales_amount")-col("prev_day_sales")))
res.show()

df_sales.createOrReplaceTempView("sales")
spark.sql("""
select
    sale_id, rep_name, date, sales_amount,
    lead(sales_amount) over( partition by rep_name order by date) as next_day_sales,
    lag(sales_amount) over(partition by rep_name order by date) as prev_day_sales,
    ( next_day_sales - sales_amount) as sales_diff_next_day,
    ( sales_amount-prev_day_sales) as sales_diff_prev_day
from sales
""").show()
```

*#NOTE: scenario 11: handling data skewness with Broadcast join*  
*# Generate transactions data using Spark directly*  
*"""*

+-----+-----+-----+			
transaction_id	customer_id	amount	
+-----+-----+-----+			
	0	cust77	84.87
	1	cust76	72.29
	2	cust63	67.45
	3	cust42	48.53
	4	cust34	47.76
	5	cust0	0.96
	6	cust52	73.6
	7	cust96	45.53
	8	cust4	3.16
+-----+-----+-----+			

only showing top 20 rows

+-----+-----+-----+		
customer_id	city	age
+-----+-----+-----+		
	cust0	city0  20
	cust1	city1  21
	cust2	city2  22
	cust3	city3  23
	cust4	city4  24
	cust5	city0  25
	cust9	city4  29
+-----+-----+-----+		

only showing top 20 rows

customer_id	transaction_id	amount	city	age
cust77	0	84.87	city2	47
cust76	1	72.29	city1	46
cust63	2	67.45	city3	33
cust42	3	48.53	city2	62
cust34	4	47.76	city4	54
cust0	5	0.96	city0	20
cust52	6	73.6	city2	22
cust96	7	45.53	city1	66

only showing top 20 rows

"""

```
transaction_df = (
    spark.range(1000000)
    .withColumn("customer_id", (100 * rand()).cast("int"))
    .withColumn("amount", round(100 * rand(), 2))
    .selectExpr("id as transaction_id",
                "concat('cust', customer_id) as customer_id", "amount")
)

# Generate customer demographic data with 100 unique customers
customer_data = [(f"cust{j}", f"city{j % 5}", 20 + j % 50) for j in range(100)]
customer_df = spark.createDataFrame(customer_data, ["customer_id", "city", "age"])

transaction_df.show()
customer_df.show()
broadcast_joined_df = (transaction_df
    .join(broadcast(customer_df), "customer_id", "inner"))
broadcast_joined_df.show()
transaction_df.createOrReplaceTempView("trans")
customer_df.createOrReplaceTempView("cust")
spark.sql("""
    select /*+ BROADCAST(c) */
        t.customer_id, t.transaction_id, t.amount, c.city, c.age
    from trans t
    join cust c
        on t.customer_id = c.customer_id
""").show()
```

*#note scenario 12*

*#note: complex aggregations with windowing functions*

*# calculate the moving average of sales for each product over the last 30 days*

*# and rank the top 5 customers by their total purchase value over the same period.*

"""

id	product_id	transaction_id	price	customer_id	date	moving_avg
17407	prod1	trans17407	499.1	cust16	2023-01-01	499.1
20723	prod1	trans20723	231.37	cust3	2023-01-01	365.235
25783	prod1	trans25783	458.32	cust4	2023-01-01	396.2633333333333
47100	prod1	trans47100	25.72	cust90	2023-01-01	303.6275
57201	prod1	trans57201	220.04	cust86	2023-01-01	286.90999999999997
43268	prod1	trans43268	106.56	cust93	2023-01-02	256.85166666666663
27820	prod1	trans27820	417.18	cust81	2023-01-03	279.7557142857143

o/p df required

```
+-----+-----+-----+
|customer_id|    total_purchase|rank|
+-----+-----+-----+
|    cust13|274669.94000000035|  1|
|    cust86| 273267.0899999998|  2|
|    cust65|273192.32000000024|  3|
|    cust51|270988.78000000026|  4|
|    cust72|270844.58999999997|  5|
+-----+-----+-----+
"""
```

```
# Create DataFrame
sales_df = spark.read.format("csv").options(header=True).load("output_file.csv")
sales_df.show()
print(sales_df.dtypes)
sales_df = sales_df.withColumn("price", round(sales_df["price"], 2))
sales_df = sales_df.withColumn("date", to_date("date"))
# sales_df = sales_df.withColumn()
print(sales_df.dtypes)

# Show sample data

window_spec = Window.partitionBy("product_id").orderBy("date").rowsBetween(-30,0)
mov_avg_30_days = sales_df.withColumn("moving_avg", avg("price").over(window_spec))

mov_avg_30_days.show()

window_spec_customer = Window.orderBy(col("total_purchase").desc())

# Group by customer and calculate total purchase in the last 30 days
customer_purchase_df = (sales_df
    .groupBy("customer_id")
    .agg(sum("price").alias("total_purchase"))
    .withColumn("rank", dense_rank().over(window_spec_customer))
)

#
top_5_customers = customer_purchase_df.filter(col("rank") <= 5)
top_5_customers.show()

sales_df.createOrReplaceTempView("sales")
spark.sql("""
    with ranked_sales as (
        select
            customer_id,
            sum(price) as total_purchases,
            dense_rank() over (order by sum(price) desc) as rank
        from
            sales
        group by
            customer_id
    )
    select *
    from ranked_sales
    where rank <= 5
""").show()
```

```

# note: below is the same query as above but we are specifying the columns which are not part of the
# aggregation
# we would get below error
# [MISSING_AGGREGATION] The non-aggregating expression "id" is based on columns which are not
# participating in the GROUP BY clause.
# with ranked_sales as (
# select
# id,product_id, transaction_id, price, customer_id, date,
# sum(price) as total_purchases,
# dense_rank() over(order by sum(price) desc) as rank
# from sales
# group by customer_id
# )
# select *
# from ranked_sales
# where rank<=5

# NOTE: scenario 13: sample json structure
"""
[
  {
    "employee_id": 101,
    "name": "John Doe",
    "department": "Engineering",
    "details": {
      "age": 29,
      "address": {
        "street": "123 Elm St",
        "city": "Springfield",
        "state": "IL",
        "postal_code": "62704"
      },
      "skills": [
        {"name": "Python", "level": "Expert", "experience": 5},
        {"name": "Spark", "level": "Intermediate", "experience": 3}
      ]
    },
    "salary": 75000.50,
    "projects": [
      {"name": "Project Alpha", "budget": 100000, "duration_months": 12},
      {"name": "Project Beta", "budget": 50000, "duration_months": 6}
    ]
  },
]
"""

# info: flatten the data
# calculate the average salary by the department
# filter the employees with skill_exp>=5 and skill_level='Expert'
# filter the projects with project_budget>60000 and project_duration>6
# add a column as salary_bane when salary>=100000 then 'High' when salary>=80000 then 'Medium'
# else 'Low' # find the avg_project_duration(months) # rank the employees by salary and then by
# department
# find the total experience include the skill_experience and sum them and
# filter the employees with more than 10 years experience
"""

```

root

```
|-- department: string (nullable = true)
|-- details: struct (nullable = true)
|   |-- address: struct (nullable = true)
|   |   |-- city: string (nullable = true)
|   |   |-- postal_code: string (nullable = true)
|   |   |-- state: string (nullable = true)
|   |   |-- street: string (nullable = true)
|   |-- age: long (nullable = true)
|   |-- skills: array (nullable = true)
|   |   |-- element: struct (containsNull = true)
|   |   |   |-- experience: long (nullable = true)
|   |   |   |-- level: string (nullable = true)
|   |   |   |-- name: string (nullable = true)
|-- employee_id: long (nullable = true)
|-- name: string (nullable = true)
|-- projects: array (nullable = true)
|   |-- element: struct (containsNull = true)
|   |   |-- budget: long (nullable = true)
|   |   |-- duration_months: long (nullable = true)
|   |   |-- name: string (nullable = true)
|-- salary: double (nullable = true)
"""

df = spark.read.option("multiline", "true").json("gpt2.json")
df.show()
df.printSchema()
trr=(df
    .withColumn("skills", expr("explode(details.skills)"))
    .withColumn("projects", expr("explode(projects)"))
)
exploded_df = trr.selectExpr(
    "employee_id",
    "name",
    "department",
    "salary",
    "skills.level as skill_level",
    "skills.experience as skill_exp",
    "skills.name as skill_name",
    "projects.budget as project_budget",
    "projects.duration_months as project_duration",
    "projects.name as project_name",
)
exploded_df.show()
exploded_df.printSchema()

avg_salary = exploded_df.groupBy("department").avg("salary")
print("average salary df")
avg_salary.show()

sep_df = exploded_df.select(
    "project_name",
    "project_duration",
    "project_budget",
    "skill_name",
    "skill_level",
    "skill_exp"
)

sep_df.show()

fil_com_df = exploded_df.filter("skill_exp>=5 and skill_level='Expert'")
print("filter exp & salary df")
fil_com_df.show()

fil_com_df_2 = exploded_df.filter("project_budget>60000 and project_duration>6")
print("filter duration>6 & budget > 60k df")
fil_com_df_2.show()
```

```

high_low_df = exploded_df.withColumn(
    "salary_band",
    expr("""
case
    when salary>=100000 then 'High'
    when salary>=80000 then 'Medium'
    else 'Low'
end
""")
)
print("salary band")
high_low_df.show()
total_budget_df = exploded_df.groupBy("department").agg(sum("project_budget").alias("total_budget"))
print("total_budget df")
total_budget_df.show()

avg_project_duration = (exploded_df
    .groupBy("department")
    .agg(avg("project_duration").alias("avg_project_duration(months)")))
print("average project duration")
avg_project_duration.show()

window_spec = Window.partitionBy("department").orderBy("salary")
rank_emp_by_dept = (exploded_df
    .withColumn("rank_by_dept", dense_rank().over(window_spec))
)
print("employee rank by the department")
rank_emp_by_dept.show()

window_spec = Window.orderBy(col("salary").desc())

salary_rank_df = (exploded_df
    .withColumn("salary_rank", dense_rank().over(window_spec)))
print("over all salary rank")
salary_rank_df.show()

total_years_df = exploded_df.groupBy("name").agg( sum("skill_exp").alias("total_exp") )
print("total+skill experience")
total_years_df.show()

max_exp_df = (exploded_df
    .groupBy("name")
    .agg(sum("skill_exp").alias("total_exp"))
    .orderBy(col("total_exp").desc())
    .filter("total_exp>10")
)
print("df with exp > 10")
max_exp_df.show()

exploded_df.createOrReplaceTempView("exploded")
spark.sql("""
select
    department,
    avg(salary) as avg_salary
from
    exploded
group by
    department
order by
    department
""").show()

```



```
# # fil_com_df = exploded_df.filter("skill_exp>=5 and skill_level='Expert'")
```

```
spark.sql("""
    select
        employee_id, name, department, salary, skill_level,
        skill_exp, skill_name, project_budget, project_duration,
        project_name
    from
        exploded
    where
        skill_exp>=5
        and
        skill_level='Expert'
""").show()
```

```
spark.sql("""
    select
        employee_id, name, department, salary, skill_level,
        skill_exp, skill_name, project_budget, project_duration,
        project_name
    from
        exploded
    where
        project_duration>6
        and
        project_budget>60000
""").show()
```

```
spark.sql("""
    select
        employee_id, name, department, salary, skill_level,
        skill_exp, skill_name, project_budget, project_duration,
        project_name,
        case
            when salary >=100000 then 'High'
            when salary>=80000 then 'Medium'
            else 'Low'
        end as salary_band
    from
        exploded
""").show()
```

```
spark.sql("""
    select
        department,
        sum(project_budget) total_budget
    from
        exploded
    group by
        department
""").show()
```

```
spark.sql("""
    select
        department,
        avg(project_duration) as avg_project_duration_months
    from
        exploded
    group by
        department
""").show()
```

```
spark.sql("""
    select
        employee_id, name, department, salary, skill_level,
        skill_exp, skill_name, project_budget, project_duration,
        project_name,
        dense_rank()over(partition by department order by salary desc) as rank_by_dept
    from
        exploded
""").show()
```

```
spark.sql("""
    select
        employee_id, name, department, salary, skill_level,
        skill_exp, skill_name, project_budget, project_duration,
        project_name,
        dense_rank()over(order by salary desc) as rank_by_salary
    from
        exploded
""").show()
```

# the same query but if you want second or third highest filter the rank for that

```
spark.sql("""
    select
        name,
        sum(skill_exp) as total_exp
    from
        exploded
    group by
        name
    having
        sum(skill_exp)>10
""").show()
```

=====

*#note: GPT-2 data*

*#note scenario 14*

```
df = spark.read.option("multiline","true").json("gpt_data.json")
```

```
# df.show()
```

```
# df.printSchema()
```

```
df = df.withColumn("products", expr("explode(products)"))
```

```
# cols = df.select("products.shipping.*")
```

```
# print(cols.columns)
```

```

exploded_df = df.selectExpr(
    "comments",
    "order_date",
    "order_id",
    "order_total",
    "products.category",
    "products.discount",
    "products.name as product_name",
    "products.price",
    "products.product_id",
    "products.quantity",
    "products.shipping.company",
    "products.shipping.estimated_delivery",
    "products.shipping.shipping_cost",
    "products.shipping.tracking_number",
    "shipped",
    "user.account_balance",
    "user.email",
    "user.loyalty_points",
    "user.name as user_name",
    "user.payment_method",
    "user.user_id",
    "user.address.city",
    "user.address.street_address",
    "user.address.zip_code",
    "user.address.geo.lat",
    "user.address.geo.lon",
)
exploded_df.printSchema()
print("EXPLODED_DF PRINTING")
exploded_df.show()

# note: 15. Customer Segmentation and Loyalty Analysis
# Scenario: Group customers based on their total spending (order_total) and loyalty points
# (loyalty_points). Create custom segments like:
# High-value, high-loyalty (e.g., customers with order totals > $5000 and loyalty points > 8000).
# Low-value, low-loyalty (e.g., order totals < $500 and loyalty points < 2000).
# Objective: Write Spark code to identify and categorize customers into these segments.
# Additionally, find the top 5 cities with the highest number of high-value customers.

loyalty_value_df = exploded_df.withColumn("Value", expr("""
    case
        when order_total>3500 then 'High'
        when order_total<500 then 'Low'
        else 'Moderate'
    end
""")).withColumn("Loyalty", expr("""
    case
        when loyalty_points>8000 then 'High'
        when loyalty_points<2000 then 'Low'
        else 'Moderate'
    end
"""))
loyalty_value_df.show()
High_value_df = loyalty_value_df.filter("Value='High'")
High_value_df.show()
# to find the top ranking cities first assign high-low-moderate to the columns based on that
# filter out the High value rows
# group them by the city and then dense rank them which will be the expected result
window_spec = Window.orderBy(col("count").desc())
high_value_df = (
    High_value_df.groupBy("city").agg(count('*').alias("count"))
    .withColumn("rank", dense_rank().over(window_spec))
    .filter("rank<=5")
)
high_value_df.show()

```

```

exploded_df.createOrReplaceTempView("exploded")
high_value_sql=spark.sql("""
    with ranking as (
        select
            *,
            case
                when order_total>3500 then 'High'
                when order_total<500 then 'Low'
                else 'Moderate'
            end as Value,
            case
                when loyalty_points>8000 then 'High'
                when loyalty_points<2000 then 'Low'
                else 'Moderate'
            end as Loyalty
        from
            exploded
    )
    select *
    from
        ranking
    where
        Value='High'
""")
high_value_sql.show()
high_value_sql.createOrReplaceTempView("high_value")
spark.sql("""
    with ranking as(
        select
            city,
            count(*) as count,
            dense_rank() over(order by count(*) desc) as rank
        from
            high_value
        group by
            city
    )
    select
        *
    from
        ranking
    where
        rank<=5
""").show()

```

```

# # note 16: below is the attempted solution for which i had to do the groupBy which i was not
# able to perform
# #INFO: keep in mind that when ever you want to perform aggregations on even windows you
# should do the groupby
# over the columns of any sort
# # or else it would not work grouping is mandatory
#
high_value_df = (
    High_value_df
        .withColumn("count", count("*").over(Window.partitionBy("city")))
        .withColumn("Cities_rank", dense_rank().over(window_spec))
        .filter("Cities_rank <= 5")
)
high_value_df.show()

```

```

#note:17. Time-Series Analysis for Shipping Efficiency
# Scenario: Calculate the difference between order_date and estimated_delivery for all orders and
# categorize them based on shipping time
# (e.g., "fast" for less than 5 days, "normal" for 5-10 days, and "slow" for more than 10 days).
# Objective: Analyze the shipping performance by shipping company and identify which company
# provides the fastest delivery on average.
# Also, calculate the average shipping cost per company.
#
raw_df = exploded_df.selectExpr(
    "order_date",
    "estimated_delivery",
    "company",
    "shipping_cost",
)
raw_df.show()
raw_df.printSchema()
days_df = raw_df.withColumn("order_date", to_timestamp(col("order_date"))) \
    .withColumn("estimated_delivery", to_timestamp(col("estimated_delivery"))) \
    .withColumn(
        "days_difference",
        datediff(col("estimated_delivery"), col("order_date"))) \
    .withColumn("delivery_category", expr("""
        case
            when days_difference > 200 then 'slow'
            when days_difference > 100 then 'normal'
            when days_difference > 50 then 'fast'
            when days_difference > 5 then 'super_fast'
            else 'within four days'
        end
        """))
days_df.show()
print("the days printed schema")
days_df.printSchema()
num_orders_df = days_df.groupBy("company").agg(count("*").alias("num_orders")).filter("num_orders > 20")
num_orders_df.show()
joined_df = days_df.join(num_orders_df, "company", "inner")
print("JOINED DF")
joined_df.show()

window_spec = Window.orderBy("avg_days_to_ship")
company_df = (
    joined_df
        .groupBy("company")
        .agg(avg("days_difference").alias("avg_days_to_ship"))
        .withColumn("rank_by_avg_delivery_time(days)", dense_rank().over(window_spec))
)
company_df.show()
print("company df showing now below\n\n\n")

avg_shipping_cost_df = (days_df
    .groupBy("company")
    .agg(avg("shipping_cost").alias("avg_shipping_cost_company"))
    .orderBy("avg_shipping_cost_company"))
avg_shipping_cost_df.show()

total_df = company_df.join(avg_shipping_cost_df, "company", 'inner')
total_df.show()
exploded_df.createOrReplaceTempView("raw")

```

```

order = spark.sql("""
    with date_conversion as (
        select
            to_timestamp(replace(order_date, 'T', ' '), 'yyyy-MM-dd HH:mm:ss') AS order_date,
            to_timestamp(estimated_delivery, 'yyyy-MM-dd') AS estimated_delivery,
            company,
            shipping_cost
        from
            raw
    ),
    days_diff as (
        select
            *,
            datediff(estimated_delivery, order_date) as days_to_deliver
        from
            date_conversion
    )
select
    *,
    case
        when days_to_deliver > 200 then 'slow'
        when days_to_deliver > 100 then 'normal'
        when days_to_deliver > 50 then 'fast'
        when days_to_deliver > 5 then 'super_fast'
        else 'within four days'
    end as order_speed
from days_diff
""")
order.show()
print("Joined data frame")

```

```

order.createOrReplaceTempView("order")

```

```

min_orders=spark.sql("""
    with company as (
        select
            company,
            count(*) as num_orders
        from
            order
        group by
            company
    )
select
    *
from
    company
where
    num_orders >= 20
""")

```

```

min_orders.createOrReplaceTempView("min_orders")

```

```

joined= spark.sql("""
    select
        o.*,
        m.num_orders
    from
        order o
    inner join
        min_orders m
    on
        m.company=o.company
""")
joined.createOrReplaceTempView("joined")

```

```

company = spark.sql("""
    with avg as (
        select
            company,
            avg(days_to_deliver) as avg_days_to_deliver
        from
            joined
        group by
            company
    )
    select
        a.*,
        dense_rank()over(order by avg_days_to_deliver) as rank_by_avg_to_ship_days
    from
        avg a
""")
company.show()

```

*#note: 18. Product Sales and Discount Optimization*

*# Scenario: Analyze product sales by category, taking into account the discounts applied (discount), and assess whether higher discounts lead to higher sales.*

*# Objective: Write a Spark query to find the correlation between discount and quantity sold.*

*# Find the top 3 product categories where higher discounts drive significantly higher sales. Additionally, analyze which product category provides the highest revenue (considering both price and quantity).*

```

raw_df = exploded_df.selectExpr(
    "order_id",
    "category",
    "discount",
    "product_name",
    "price",
    "quantity",
    "order_total"
)
raw_df.show()
raw_df.printSchema()

correlation_df = (raw_df
    .groupBy("category")
    .agg(corr("discount", "quantity").alias("correlation")))

correlation_df.show()

top_3_categories = correlation_df.orderBy(col("correlation").desc()).limit(3)
top_3_categories.show()
revenue_df = (raw_df
    .groupBy("category")
    .agg(sum(col("price") * col("quantity")).alias("total_revenue")))
top_revenue_category = revenue_df.orderBy(col("total_revenue").desc())
top_revenue_category.show()

# Assuming your DataFrame is named 'df'
average_order_total = raw_df.agg(avg("order_total"))
average_order_total.show()

```

```

raw_df.createOrReplaceTempView("raw")
spark.sql("""
    select
        category,
        sum(price * quantity) as total_revenue
    from
        raw
    group by
        category
    order by
        total_revenue desc
""").show()

```

```

spark.sql("""
    select
        avg(order_total) as avg_order_price
    from
        raw
""").show()

```

*#note: 19. Fraud Detection with Anomalous Orders*

*# Scenario: Identify potentially fraudulent transactions where the account\_balance is suspiciously low (e.g., < \$50), but the order\_total is abnormally high (e.g., > \$1000), especially when the payment\_method is unconventional (e.g., bitcoin). Objective: Write a Spark query to detect such anomalies, and analyze trends (e.g., are certain product categories more susceptible to these transactions?).*

```

anomaly_df = (exploded_df
    .filter("account_balance<50 and order_total>1000 and payment_method='bitcoin'")
    .select("account_balance", "order_total", "category", "payment_method"))

```

```

anomaly_df.show()

```

```

anomaly_df = (anomaly_df
    .groupBy("category")
    .agg(count("category").alias("count"))
    .orderBy(col("count").desc()))

```

```

print(anomaly_df)

```

```

print()

```

```

print(f""""The category with got affected by frauds the most is
    {anomaly_df.collect()[0][0]} with {anomaly_df.collect()[0][1]} of frauds""")

```

```

print()

```

```

# "account_balance", "order_total", "category", "payment_method"

```

```

spark.sql("""
    select
        account_balance, order_total, category, payment_method
    from
        exploded
    where
        account_balance<50 and
        order_total>1000 and
        payment_method='bitcoin'
""").show()

```

```

spark.sql("""
    select
        category,
        count(*) as frauds
    from
        exploded
    where
        account_balance<50 and
        order_total>1000 and
        payment_method='bitcoin'
    group by
        category
    order by
        frauds desc
""").show()

```



```

# note: 20. Order Retention and Cancellation Insights
# Scenario: Analyze whether shipped (shipped = True) orders are more likely to have customer
# comments or issues, compared to non-shipped orders.
#
shipped_df = (exploded_df
              .filter("shipped=true")
              .select("comments", "shipped")
              .groupBy("comments")
              .agg(count("comments").alias("count")))
              .orderBy(col("count").desc())

shipped_df.show()
print(shipped_df.count())

spark.sql("""
select
    comments,
    count(comments) as count
from
    exploded
where
    shipped=true
group by
    comments
order by
    count desc
""").show()

```

```
# note: gpt-3 Data
# note: 21
df = spark.read.option("multiline","true").json("Ecommerce_data.json")
# df.printSchema()
df = df.withColumn("items", expr("explode(order.items)"))
df.printSchema()
```

```
# x = df.select("order.shipping.*")
# print(x.columns)
```

```
#
"""
```

```
root
|-- order: struct (nullable = true)
|   |-- customer: struct (nullable = true)
|   |   |-- account_balance: double (nullable = true)
|   |   |-- email: string (nullable = true)
|   |   |-- loyalty: struct (nullable = true)
|   |   |   |-- loyalty_points: long (nullable = true)
|   |   |   |-- tier: string (nullable = true)
|   |   |-- name: struct (nullable = true)
|   |   |   |-- first_name: string (nullable = true)
|   |   |   |-- last_name: string (nullable = true)
|   |   |-- user_id: string (nullable = true)
|   |-- items: array (nullable = true)
|   |   |-- element: struct (containsNull = true)
|   |   |   |-- category: string (nullable = true)
|   |   |   |-- name: string (nullable = true)
|   |   |   |-- price: double (nullable = true)
|   |   |   |-- product_id: string (nullable = true)
|   |   |   |-- quantity: long (nullable = true)
|   |-- order_date: string (nullable = true)
|   |-- order_id: string (nullable = true)
|   |-- payment: struct (nullable = true)
|   |   |-- amount: double (nullable = true)
|   |   |-- discount: double (nullable = true)
|   |   |-- method: string (nullable = true)
|   |-- shipping: struct (nullable = true)
|   |   |-- address: struct (nullable = true)
|   |   |   |-- city: string (nullable = true)
|   |   |   |-- coordinates: struct (nullable = true)
|   |   |   |   |-- lat: double (nullable = true)
|   |   |   |   |-- lon: double (nullable = true)
|   |   |   |-- street_address: string (nullable = true)
|   |   |   |-- zip_code: string (nullable = true)
|   |   |-- estimated_delivery: string (nullable = true)
|   |   |-- shipped: string (nullable = true)
|   |   |-- shipping_cost: double (nullable = true)
|   |   |-- tracking_number: string (nullable = true)
|-- items: struct (nullable = true)
|   |-- category: string (nullable = true)
|   |-- name: string (nullable = true)
|   |-- price: double (nullable = true)
|   |-- product_id: string (nullable = true)
|   |-- quantity: long (nullable = true)
```

```
"""
```

```

exploded_df = df.selectExpr(
    "order.customer.account_balance",
    "order.customer.email",
    "order.customer.loyalty.loyalty_points",
    "order.customer.loyalty.tier",
    "order.customer.name.first_name",
    "order.customer.name.last_name",
    "order.customer.user_id",
    "items.category",
    "items.name",
    "items.price",
    "items.product_id",
    "items.quantity",
    "order.order_date",
    "order.order_id",
    "order.payment.amount",
    "order.payment.discount",
    "order.payment.method",
    "order.shipping.address.city",
    "order.shipping.address.coordinates.lat",
    "order.shipping.address.coordinates.lon",
    "order.shipping.address.street_address",
    "order.shipping.address.zip_code",
    "order.shipping.estimated_delivery",
    "order.shipping.shipped",
    "order.shipping.shipping_cost",
    "order.shipping.tracking_number",
)

# exploded_df.show()
exploded_df.printSchema()
#note: 1. Calculate the Average Shipping Time per User Tier
# You need to calculate the average number of days it takes for an order to be delivered for each
# user based on their tier (standard, premium, etc.).
# Use the order_date and estimated_delivery columns to calculate the number of days for delivery,
# and then aggregate the results by tier.
# Transformation Steps:
# Calculate the difference between estimated_delivery and order_date to get days_for_delivery.
# Group by tier and calculate the average delivery time for each tier.
# Filter orders that have been marked as shipped = true.
exploded_df = (exploded_df
    .withColumn(
        "order_date",
        to_date(exploded_df["order_date"], "yy-MM-dd")
    )
    .withColumn(
        "estimated_delivery",
        to_date(exploded_df["estimated_delivery"], "yy-MM-dd")
    )
    .withColumn(
        "days_for_delivery",
        date_diff(col('estimated_delivery'), col('order_date'))
    )
)

# exploded_df = exploded_df.filter("days_for_delivery>0")
exploded_df = exploded_df.withColumn("rating", (rand()*5).cast("int"))
exploded_df.printSchema()
# exploded_df.cache()
exploded_df.show()
exploded_df.createOrReplaceTempView("exploded")

raw_df = (exploded_df
    .filter(exploded_df["shipped"]=="true")
    .groupBy("tier")
    .agg(avg("days_for_delivery").alias("avg_days_for_delivery_by_tier"))
)
raw_df.show()

```

```

spark.sql("""
    select
        tier,
        avg(days_for_delivery) as avg_days_to_ship_by_tier
    from
        exploded
    where
        shipped='true'
    group by
        tier
""").show()

```

*#note: 2. Join Orders with User Information and Perform Aggregations*

*# Create a new DataFrame by joining the user's profile information with their orders.*

*# Then, compute the following metrics: #*

*# Total order\_amount for each user.*

*# Total loyalty\_points gained for each purchase.*

*# Average discount received per user across all their orders.*

*# Transformation Steps:*

*#*

*# Perform an inner join between the user information (using user\_id) and the order details.*

*# Create new columns for total order amount (price \* quantity - discount).*

*# Aggregate data at the user level to compute total order\_amount, loyalty\_points, and average discount.*

```

raw_df = exploded_df.selectExpr(
    "user_id",
    "account_balance",
    "category",
    "loyalty_points",
    "order_id",
    "product_id",
    "quantity",
    "discount",
    "price",
    "shipped",
    "tier"
)
raw_df.printSchema()
cust_df = (raw_df
    .groupBy("user_id")
    .agg(
        avg("discount").alias("avg_discount"),
        sum(col("price")*col("quantity")-col("discount")).alias("total_order_amount")
    )
)
cust_df.show()
tier_df = raw_df.groupBy("tier").agg(avg("discount").alias("avg_discount_by_tier"))
tier_df.show()
raw_df.createOrReplaceTempView("raw")
spark.sql("""
    select
        user_id,
        avg(discount)as avg_discount,
        sum(price*quantity-discount) as total_order_amount
    from
        raw
    group by
        user_id
""").show()

```

```
spark.sql("""
select
    tier,
    avg(discount) as avg_discount_by_tier
from
    raw
group by
    tier
""").show()
```

```
#note: 3. Identify Users Who Are Eligible for Tier Upgrade Based on Spending Patterns
# Assume that users are eligible for a tier upgrade if their total order amount across all
# orders exceeds a
# threshold (e.g., $5,000). Generate a list of users who are eligible for a tier upgrade
# and mark their tier as premium.
#
# Transformation Steps:
#
# Aggregate orders by user_id to calculate total spending.
# Filter users whose total spending exceeds the threshold.
# Update the tier column in the resulting DataFrame to reflect the new tier status for eligible users.
```

```
raw_df = exploded_df.selectExpr(
    "user_id",
    "account_balance",
    "category",
    "loyalty_points",
    "order_id",
    "product_id",
    "quantity",
    "discount",
    "price",
    "shipped",
    "tier"
)
cust_df = (raw_df
    .groupBy("user_id")
    .agg(sum(col("price")*col("quantity")-col("discount")).alias("total_order_amount")))
cust_df = cust_df.join(raw_df.select("user_id", "tier").distinct(), on="user_id", how="left")
```

```
# Step 3: Update the 'tier' based on the condition
# cust_df = cust_df.withColumn(
#     "tier", when(col("total_order_amount") > 5000, "Premium").otherwise(col("tier"))
# )
cust_df = cust_df.withColumn("tier", expr("""
    case
        when total_order_amount>5000 then 'Premium'
        else tier
    end
"""))
cust_df.show()
```

```
# Extract the user_id column and collect the values into a list
Premium_user_id_list = cust_df.filter("tier='Premium'").select("user_id").rdd.flatMap(lambda x: x).collect()

# Show the list of user_id values
print(len(Premium_user_id_list))
raw_df.createOrReplaceTempView("raw")
```

#### *#note: 4. Complex Nested JSON Scenario*

```
# Assume you receive new data in JSON format with nested fields, containing user reviews for products.
# The JSON schema is as follows:
# Flatten the nested structure, extracting user_id, order_id, product_id, rating, and comments.
# Aggregate the average rating per product_id and identify the products with the highest average rating.
# Transformation Steps:
# Group by product_id to compute the average rating and rank products based on ratings.
```

```
raw_df = exploded_df.selectExpr("product_id", "user_id", "rating")
raw_df = (raw_df
          .groupBy("product_id")
          .agg(avg("rating").alias("avg_rating"))
          .orderBy(col("avg_rating").desc())
        )
raw_df.show()
```

#### *# note: 5. Identify Top Performing Cities by Revenue*

```
# Calculate the total revenue generated from orders for each city and then rank the cities
# based on their revenue.
# Only include orders where the product was shipped (shipped = true). Transformation Steps:
# Filter orders where shipped = true.
# Compute total revenue per city (quantity * price - discount).
# Rank cities based on total revenue and return the top 10.
```

```
raw_df = exploded_df.filter("shipped='true'")
window_spec = Window.orderBy(col("total_order_amount_by_city").desc())
cities_df = (raw_df
             .withColumn("total_order_amount", expr("price*quantity-discount"))
             .groupBy("city")
             .agg(sum("total_order_amount").alias("total_order_amount_by_city"))
             .withColumn("rank",dense_rank().over(window_spec))
            )
cities_df.show()
```